

# KUBERNETES

Basics

Gaurav Sharma


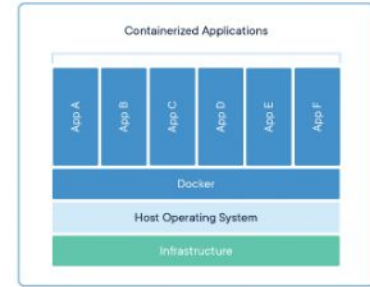


# AGENDA

- 👉 What is a Container ?
- 👉 What is Docker?
- 👉 Traditional Vs Containerized Deployment
- 👉 What is Kubernetes and how it works ?
- 👉 Why you need Kubernetes ?
- 👉 Kubernetes Cluster Architecture
- 👉 Control Plane Vs Data Plane
- 👉 Master Node : Scheduler & Etcd
- 👉 Worker Node : Container Runtime Engine
- 👉 What is Pod ?
- 👉 How multiple containers are managed in a Pod?
- 👉 What is a Namespace ?
- 👉 What is a Replication Set ?
- 👉 What is a Deployment ?
- 👉 What is a Service ?
- 👉 Storage and Networking in Kubernetes
- 👉 Security in Kubernetes

# What is a container ?

- **Encapsulate applications** and their **dependencies**, including runtime, **libraries** and settings, ensuring consistency across different environments
- Containers **share** the **host OS kernel** making them lightweight and resource-efficient compared to traditional virtual machines
- Provide process isolation, securing applications from **interference** and enhancing security

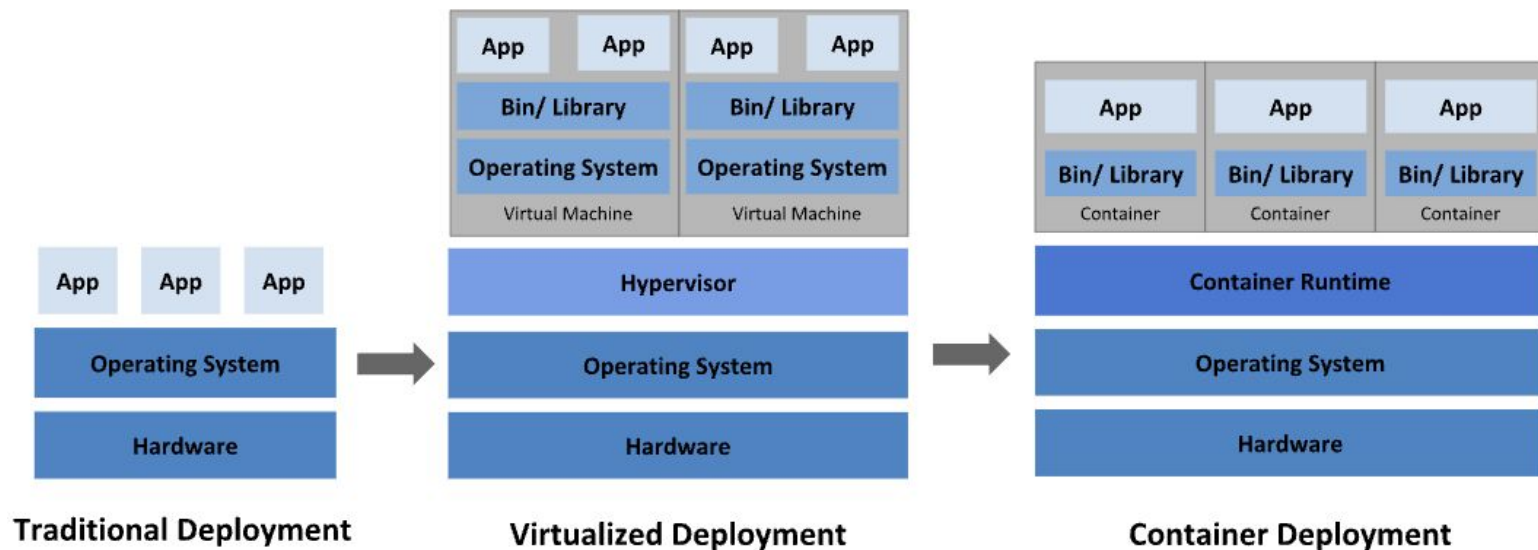


# What is Docker ?

- Docker is an open platform for **developing**, **shipping**, and **running** applications.
- It enables you to separate your applications from your infrastructure so you can deliver software quickly
- It provides the ability to package and run an application in a **loosely isolated environment** called container
- With Docker, you can manage your infrastructure in the same way as you manage your applications
- Docker official documentation [Link](#)



# Traditional vs Containerized Deployment



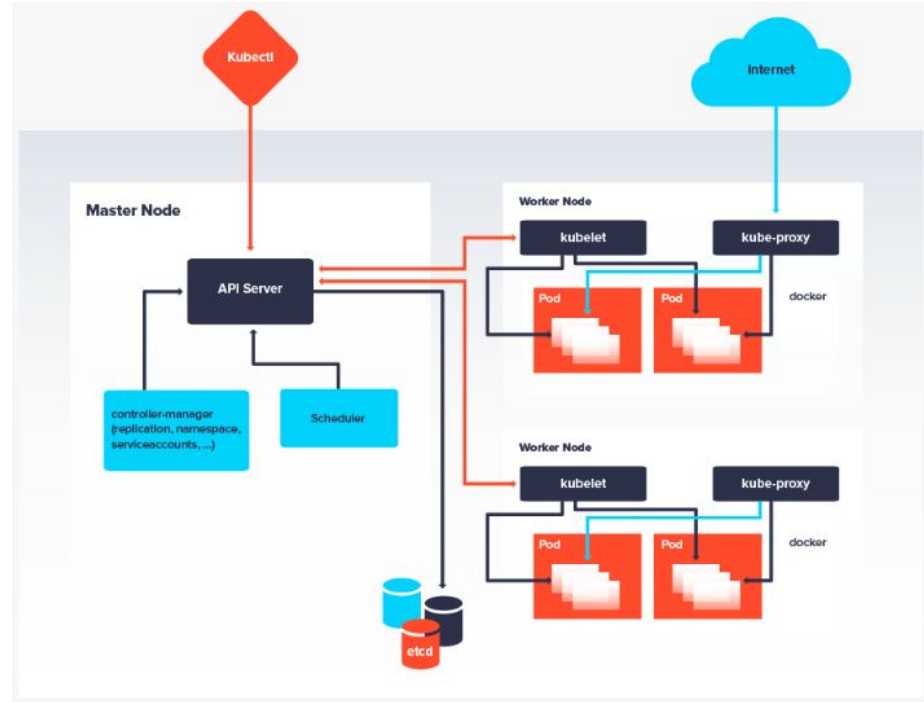
# What is Kubernetes ?

- Word “Kubernetes” is originated from greek word which means helmsman/pilot
- Also called as **K8s**,K8s as an abbreviation results from counting the eight letters between the "k" and "s"
- **Portable, extensible**, open source platform for managing **containerized** workloads\services
- **Orchestrator** that schedules containers on a cluster and manages **workloads** to ensure they run as intended
- Facilitates both declarative configuration and automation
- K8s works by managing and coordinating containers across a cluster of machines



# How Kubernetes works ?

- Master Node functions as the **central orchestrator**, **managing workloads** and overseeing critical components such as the API Server, Controller Manager, Scheduler, and etcd
- Nodes serve as the **execution environment** for applications, receiving instructions and updates from the Master.
- Etcd component functions as a **distributed key-value store**, storing the **cluster's state**.
- Kubelet ensures container health and **facilitates communication**
- Kubernetes API Server **validates requests** and maintains the cluster's integrity in a cohesive and professional manner.



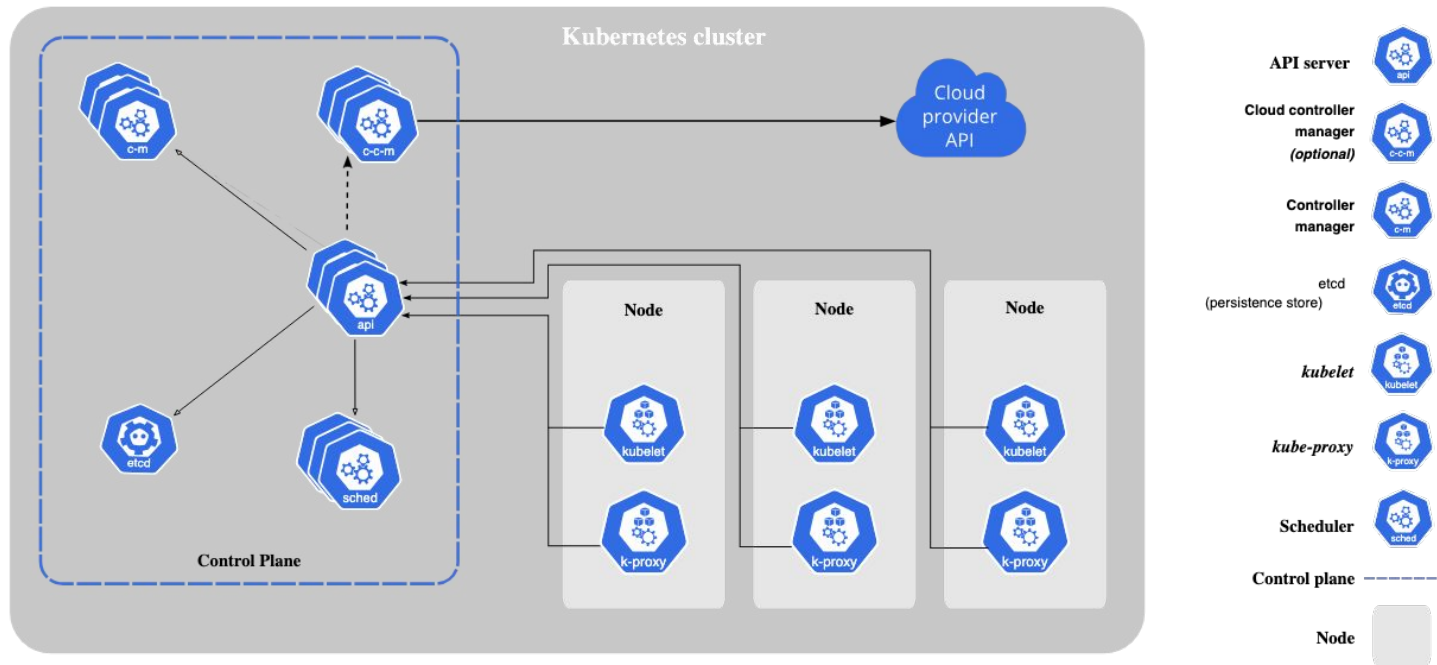
# Why you need Kubernetes?

- 1 Service discovery & load balancing
- 2 Storage orchestration
- 3 Automated rollouts and rollbacks
- 4 Automatic bin packing
- 5 Self-healing
- 6 Secret & configuration management
- 7 Horizontal scaling





# Kubernetes Cluster Architecture



# Control Plane

## kube-apiserver

- API server is the frontend for the K8s control plane that exposes the K8s API
- It is designed to scale horizontally— that is, it scales by deploying more instances.
- It validates and configures data for the api objects which include pods, services, replicationcontrollers, and others

## etcd

- Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data

## kube-scheduler

- Control plane component that watches for newly created Pods with no assigned node, and selects a node for them to run on
- Factors taken into account for scheduling decisions include: individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference, and deadlines

## kube-controller-manager

- Control plane component that runs controller processes
- Logically, each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process
- different types of controllers - Node controller, Job controller, Service account controller

## cloud-controller -manager

- A Kubernetes control plane component that embeds cloud-specific control logic
- It runs controllers that are specific to your cloud provider
- It combines several logically independent control loops into a single binary that you run as a single process



# Data Plane

## Kubelet

- An agent that runs on each node in the cluster. It makes sure that containers are running in a Pod
- It takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy
- Kubelet doesn't manage containers which were not created by Kubernetes

## Kube-Proxy


- kube-proxy is a network proxy that runs on each node in your cluster, implementing part of the Kubernetes Service concept
- It maintains network rules on nodes that allow network communication to the pods from network sessions inside or outside of the cluster
- It uses the operating system packet filtering layer if there is one and it's available

## Container-Runtime

- A fundamental component that empowers Kubernetes to run containers effectively
- It is responsible for managing the execution and lifecycle of containers within the Kubernetes environment
- Kubernetes supports container runtimes such as containerd, CRI-O, and any other implementation of the Kubernetes CRI (Container Runtime Interface)



# Master Node



# Control Plane : Scheduler

- Kube-scheduler is the **default scheduler** for Kubernetes
- The scheduler finds feasible nodes for a pod and then runs a set of functions to score the feasible nodes
- The node with the highest score is selected among the feasible ones to run the Pod and the API server is then notified. This process is called **binding**
- Kube-scheduler selects a node for the pod in a 2-step operation: **Filtering & Scoring**
  - The filtering step finds the set of Nodes where it's feasible to schedule the Pod
  - In the scoring step, the scheduler ranks the remaining nodes to choose the most suitable Pod placement and assigns a score to each node that survived filtering
- There are two supported ways to configure the filtering and scoring behavior of the scheduler :
  - **Scheduling Policies**
  - **Scheduling Profiles**

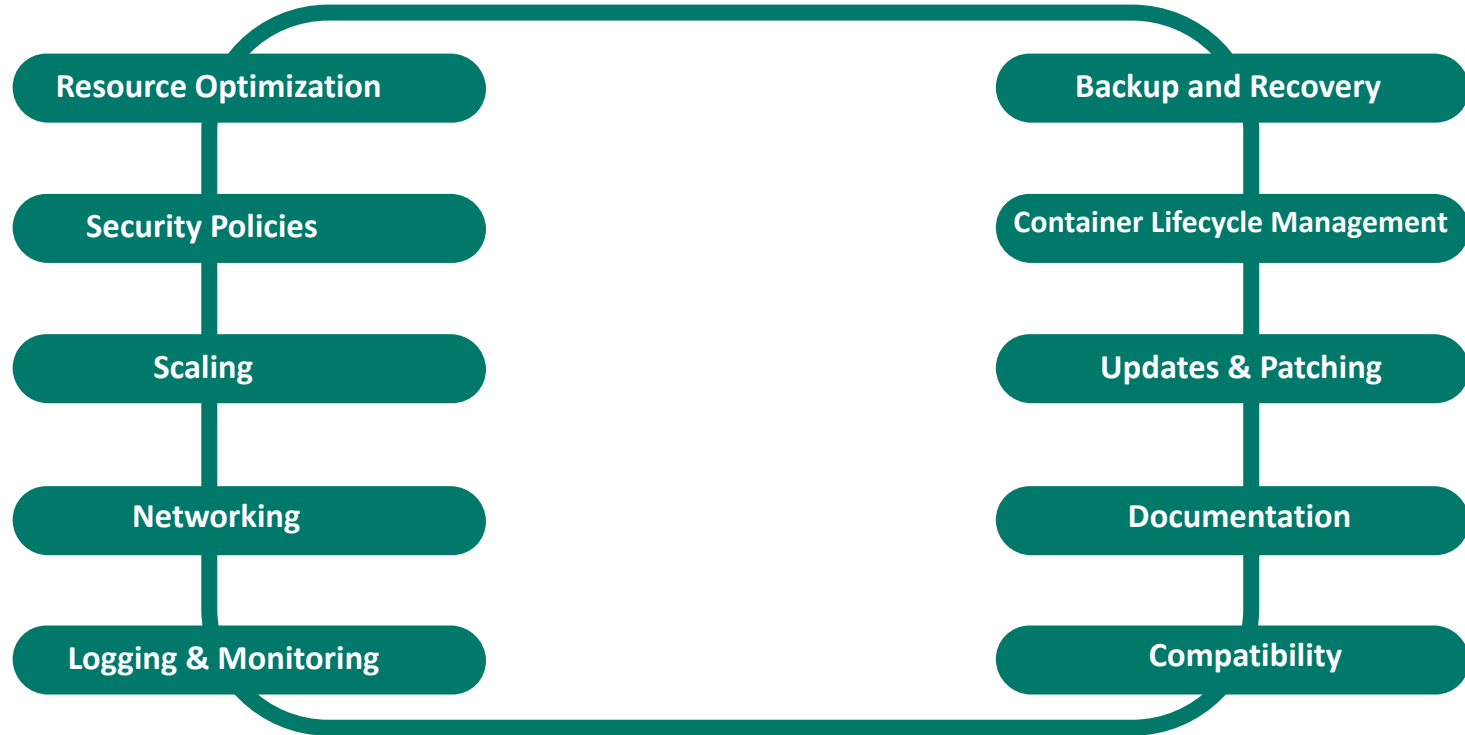


# Control Plane : Etcd

- **Distributed Key-Value Store** : Etcd is a distributed, consistent key-value store used for configuration management and service discovery in Kubernetes.
- **Consistency and Reliability** : It ensures strong consistency, allowing for reliable storage and retrieval of configuration data across a Kubernetes cluster.
- **Raft Consensus Algorithm** : Etcd employs the Raft consensus algorithm to maintain consistency and fault tolerance among distributed nodes.
- **Critical Cluster Component** : It serves as a critical component in a Kubernetes cluster, storing essential information about the cluster's state.
- **API for Kubernetes** : Etcd provides a reliable API that Kubernetes components use to store and retrieve critical configuration and state information.
- **Secure Communication** : Security is paramount, and etcd supports secure communication through encryption and access control mechanisms, ensuring the integrity of stored data.



# Worker Node



# Data Plane : Container Runtime Engine





# Pod



- **Smallest** deployable units of computing in k8s
- Group of one or more containers with shared storage and network resources
- Pods in a Kubernetes cluster are used in two main ways:
  - Pods that run a **single** container
  - Pods that run **multiple** containers that need to work together
- Pods are generally not created directly and are created using workload resources
- Containers in a Pod can **share resources**, **dependencies**, and communicate seamlessly.


Sample Pod defn yaml file

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```




# How Pods manage multiple containers ?

- **Automatic Co-location and Scheduling:** Pods streamline container management through automatic co-location on the same machine, promoting efficiency.
- **Seamless Resource and Dependency Sharing:** Containers within a Pod effortlessly share resources and dependencies, facilitating collaborative processes.
- **Coordinated Communication:** Containers can communicate and coordinate termination within a Pod, enhancing operational cohesion.
- **Init Container Efficiency:** Init containers, running before app containers by default, exemplify efficient process initiation.
- **Init Container Restart Policy Control:** Utilizing the Sidecar Containers feature gate enhanced control over restart policies for init containers, ensuring consistent operation.
- **Networking and Storage Resources:** Pods inherently offer shared networking and storage resources, further enhancing their ability to manage multiple containers effectively.



# Namespace

- Namespace provides a mechanism for **isolating** groups of resources within a single cluster
- A way to divide cluster resources between multiple users (via **resource quota**)
- Intended for use in environments with many users spread across multiple teams, or projects
- **Cannot be nested** inside one another and each Kubernetes resource can only be in one namespace



```
apiVersion: v1
kind: Namespace
metadata:
  name: <insert-namespace-name-here>
```

Sample Namespace defn yaml file



# Replication Set

- A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time
- It is linked to its Pods via the Pods "metadata.ownerReferences" field
- It identifies new pods to acquire by using its selector
- It uses pod template to create new pods
- It is recommend to use Deployments instead of directly using ReplicaSets

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # modify replicas according to your case
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
```

Sample Replicaset defn yaml file



# Deployment

- **Application Scaling** : Deployments enable effortless scaling of applications by managing the deployment and scaling of Pods.
- **Rolling Updates** : Facilitates seamless rolling updates, ensuring continuous application availability during the update process.
- **Rollback Capability** : Offers easy rollback to previous versions in case of issues with the latest deployment.
- **Declarative Configuration** : Defined using declarative YAML, allowing easy configuration and version control.
- **Load Balancing** : Automatically provides load balancing across Pods to distribute traffic evenly.
- **Self-healing** : Monitors and ensures the desired state, automatically replacing failed Pods and maintaining application availability.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Sample Deployment defn yaml file



# Services

- Abstraction to help **expose groups of Pods** over a network.
- Each Service object defines a logical set of endpoints
- The set of Pods targeted by a Service is usually determined by a selector that is defined
- Different service types :
  - ClusterIP
  - NodePort
  - LoadBalancer
  - ExternalName
- Headless service : Designed for scenarios where direct communication with individual pods is necessary

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Sample Service defn yaml file



# Storage

- Volumes

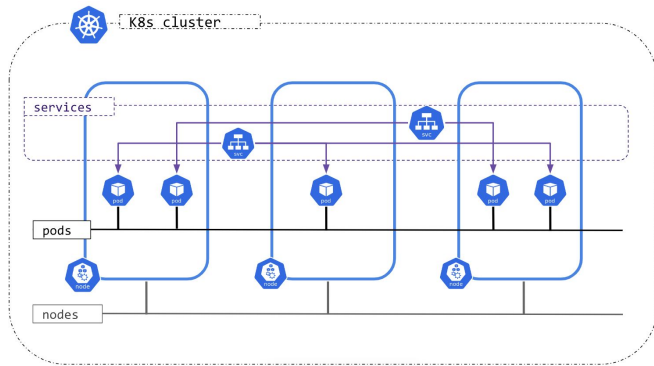
- **NFS:** Network File System volume allows pods to share files across the network. Useful for sharing data between pods.
- **ConfigMap:** Mounts configuration data as a volume, allowing pods to consume configuration files or environment variables.
- **Secret:** Similar to ConfigMap but designed for storing sensitive information such as passwords or API keys.
- **PersistentVolume (PV):** Represents physical storage in the cluster, decoupling it from individual pods. Can be dynamically provisioned.
- **PersistentVolumeClaim (PVC):** Requests a specific amount of storage from a PersistentVolume. Binds with a matching PV.

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      volumeMounts:
        - name: redis-storage
          mountPath: /data/redis
  volumes:
    - name: redis-storage
      emptyDir: {}
```

Sample Volume defn yaml file



# Networking



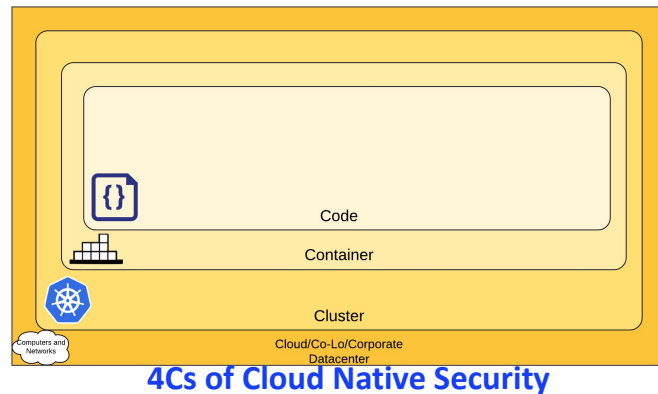
- **Pod Networking:**
  - Pods communicate with each other via a flat network within the cluster.
  - Each pod gets its unique IP address for inter-pod communication.
- **Service Networking:**
  - Services enable load balancing and provide a stable endpoint for accessing pods.
  - Cluster IP, NodePort, and LoadBalancer service types manage different networking scenarios.
- **Ingress:**
  - Manages external access to services, acting as an API gateway.
  - Routes external traffic to appropriate services based on rules and configurations.
- **Network Policies:**
  - Controls traffic between pods using defined policies.
  - Specifies which pods can communicate with each other based on labels and selectors.
- **CNI Plugins:**
  - Container Network Interface (CNI) plugins manage container networking.
  - Plugins handle tasks like IP address allocation, routing, and network isolation.
- **DNS-Based Service Discovery:**
  - Kubernetes uses DNS to enable service discovery.
  - Services and pods are assigned DNS names, allowing easy and dynamic resolution.





# Security

Security in Kubernetes is a multi-faceted approach, involving access control, network segmentation, and secure handling of sensitive information. These concepts collectively contribute to a robust security posture within a Kubernetes cluster.



- **Role-Based Access Control (RBAC)** : Ensures fine-grained access control by defining roles and role bindings. It specifies what actions users, groups, or service accounts can perform within the cluster.
- **Pod Security Policies** : Defines security policies for pods, restricting privileges and access.
- **Network Policies** : Specifies how pods communicate with each other, enforcing rules on ingress and egress traffic.
- **Secrets and ConfigMaps** : Manages sensitive information and configuration data securely. Kubernetes encrypts and manages secrets, preventing unauthorized access.



# The End

Devops/SRE-DeepDive

