



## Kubernetes Comprehensive Hands-On Guide

[Click Here To Enrol To Batch-6 | DevOps & Cloud DevOps](#)

### Introduction to Kubernetes

Kubernetes is an open-source platform designed to automate deploying, scaling, and operating application containers. It was originally developed by Google and is now maintained by the Cloud Native Computing Foundation (CNCF).

### Key Concepts

1. **Cluster:** A set of nodes (machines) running Kubernetes, managed by a control plane.
2. **Node:** A single machine in the cluster, which runs containerized applications.
3. **Pod:** The smallest deployable unit in Kubernetes, which can contain one or more containers.
4. **Deployment:** A controller that manages the deployment and scaling of Pods.
5. **Service:** An abstraction that defines a logical set of Pods and a policy by which to access them.
6. **ConfigMap:** Used to manage configuration data separately from application code.
7. **Secret:** Used to manage sensitive data, such as passwords, OAuth tokens, etc.

### Kubernetes Architecture

#### Introduction

Kubernetes, often abbreviated as K8s, is an open-source container orchestration platform designed to automate the deployment, scaling, and operation of containerized applications. It is built to handle a production environment, with multiple machines and services running simultaneously. Understanding the

architecture of Kubernetes is crucial for effectively deploying and managing applications.

## Overview of Kubernetes Architecture

Kubernetes follows a client-server architecture and comprises the following key components:

- **Control Plane:** Manages the Kubernetes cluster.
- **Node Components:** Run on every node in the cluster.
- **Objects:** Persistent entities in the Kubernetes system.

## Control Plane Components

The control plane components make global decisions about the cluster (e.g., scheduling) and detect and respond to cluster events.

1. **etcd:**
  - **Role:** A consistent and highly-available key-value store used as Kubernetes' backing store for all cluster data.
  - **Importance:** Stores configuration data, representing the state of the cluster at any given point in time.
2. **kube-apiserver:**
  - **Role:** Exposes the Kubernetes API. It is the front-end for the Kubernetes control plane.
  - **Importance:** Validates and configures the data for the API objects, which include pods, services, replication controllers, and others.
3. **kube-scheduler:**
  - **Role:** Watches for newly created Pods with no assigned node and selects a node for them to run on.
  - **Importance:** Ensures Pods are placed on nodes that have sufficient resources to run them.
4. **kube-controller-manager:**
  - **Role:** Runs controller processes.
  - **Importance:** Manages different types of controllers, including the replication controller, which ensures the specified number of pod replicas are running at any one time.
5. **cloud-controller-manager:**
  - **Role:** Runs controllers specific to the cloud provider.
  - **Importance:** Manages cloud-specific control logic, allowing for Kubernetes to be cloud-agnostic.

## Node Components

These run on every node in the cluster, ensuring containers are running in a Pod.

### 1. **kubelet:**

- **Role:** An agent that runs on each node in the cluster. It ensures containers are running in a Pod.
- **Importance:** Communicates with the control plane to receive instructions and report back the status.

### 2. **kube-proxy:**

- **Role:** Maintains network rules on nodes. These network rules allow network communication to your Pods from network sessions inside or outside of your cluster.
- **Importance:** Facilitates the Kubernetes networking model.

### 3. **Container Runtime:**

- **Role:** Software responsible for running containers.
- **Importance:** Examples include Docker, containerd, and CRI-O.

## **Persistent Storage Components**

Kubernetes supports several storage mechanisms, which allow applications to store and retrieve data.

### 1. **PersistentVolumes (PV):**

- **Role:** A piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes.
- **Importance:** Provides a way for users to store data persistently.

### 2. **PersistentVolumeClaims (PVC):**

- **Role:** A request for storage by a user.
- **Importance:** Users create PVCs to request storage resources and bind to PVs.

### 3. **Storage Classes:**

- **Role:** Define the types of storage classes available in a cluster.
- **Importance:** Allow dynamic provisioning of PVs.

## **Kubernetes Objects**

These persistent entities represent the state of the cluster. Kubernetes uses these entities to represent the desired state of the cluster and change the actual state to the desired state.

### 1. **Pods:**

- **Role:** The smallest and simplest Kubernetes object. A Pod represents a set of running containers on your cluster.
- **Importance:** The fundamental unit of deployment.

### 2. **Services:**

- **Role:** An abstract way to expose an application running on a set of Pods as a network service.

- **Importance:** Allows for Pods to communicate with each other or external services.
- 3. **ReplicaSets:**
  - **Role:** Ensures a specified number of pod replicas are running at any given time.
  - **Importance:** Provides high availability.
- 4. **Deployments:**
  - **Role:** Provides declarative updates for Pods and ReplicaSets.
  - **Importance:** Manages and maintains desired states.
- 5. **ConfigMaps and Secrets:**
  - **Role:** ConfigMaps store configuration data in key-value pairs. Secrets store sensitive information.
  - **Importance:** Manage configuration and sensitive data separately from application code.

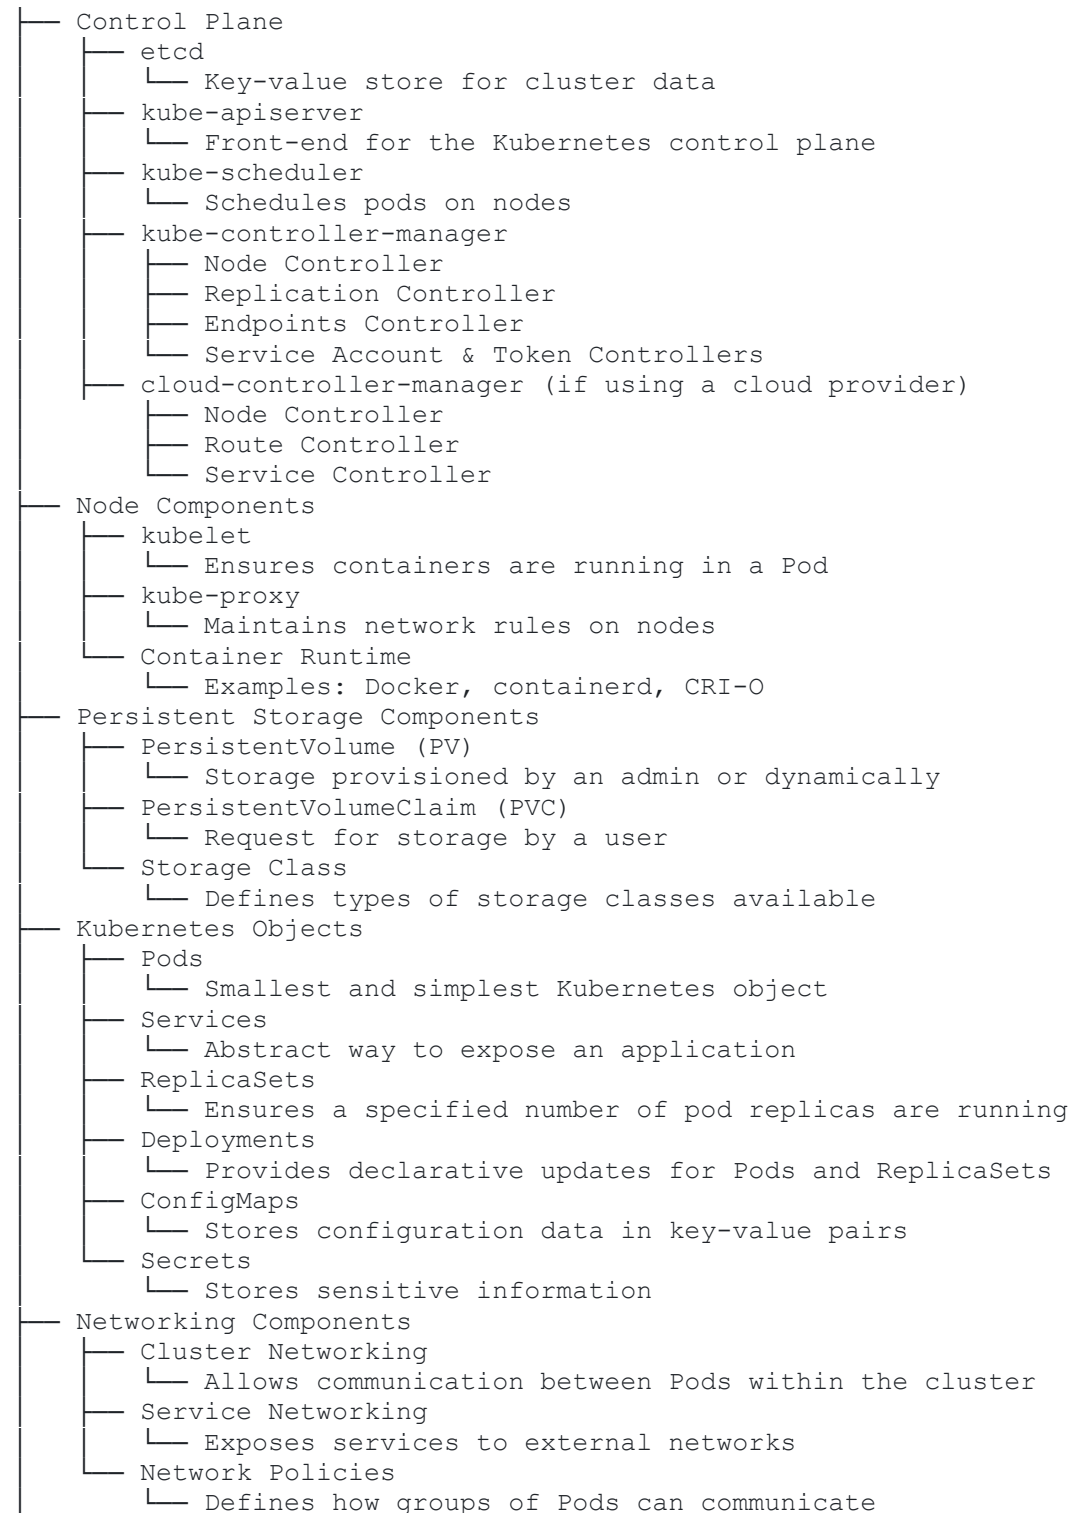
## Networking in Kubernetes

Kubernetes networking allows for communication between different components within the cluster and with the outside world.

1. **Cluster Networking:**
  - **Role:** Allows communication between Pods within the same cluster.
  - **Importance:** Ensures that each Pod can communicate with any other Pod without NAT.
2. **Service Networking:**
  - **Role:** Exposes services to external networks.
  - **Importance:** Allows external users to access services within the cluster.
3. **Network Policies:**
  - **Role:** Define how groups of Pods can communicate with each other and with other network endpoints.
  - **Importance:** Provide fine-grained control over network traffic.

# Kubernetes Architecture Diagram

## Kubernetes Cluster



## Setting Up a Kubernetes Cluster

### Minikube

Minikube is a tool that allows you to run Kubernetes locally. It's perfect for development and testing.

1. **Install Minikube:**

2. 

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64  
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

3. **Start Minikube:**

```
minikube start
```

4. **Verify Installation:**

```
kubectl get nodes
```

### Kubernetes on Cloud

You can set up a Kubernetes cluster on various cloud providers like Google Kubernetes Engine (GKE), Amazon Elastic Kubernetes Service (EKS), or Azure Kubernetes Service (AKS). Here's an example for GKE:

1. **Install Google Cloud SDK and authenticate:**

2. 

```
curl https://sdk.cloud.google.com | bash
```
3. 

```
exec -l $SHELL
```
4. 

```
gcloud init  
gcloud auth login
```

5. **Create a GKE cluster:**

```
gcloud container clusters create my-cluster --zone us-central1-a
```

6. **Get authentication credentials:**

```
gcloud container clusters get-credentials my-cluster --zone us-central1-a
```

7. **Verify cluster:**

```
kubectl get nodes
```

## Deploying Applications

### Creating a Deployment

A Deployment provides declarative updates to applications. Here's how to create one:

## 1. Create a Deployment YAML file (deployment.yaml):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
```

## 2. Apply the Deployment:

```
kubectl apply -f deployment.yaml
```

## 3. Verify the Deployment:

```
kubectl get deployments
kubectl get pods
```

### Exposing a Deployment

To expose your Deployment so that it can be accessed from outside the Kubernetes cluster, you need to create a Service.

## 1. Create a Service YAML file (service.yaml):

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

## 2. Apply the Service:

```
kubectl apply -f service.yaml
```

## 3. Verify the Service:

```
kubectl get services
```

## Scaling Applications

Scaling applications in Kubernetes is straightforward.

### 1. Scale the Deployment:

```
kubectl scale deployment nginx-deployment --replicas=5
```

### 2. Verify the Scaling:

```
kubectl get deployments  
kubectl get pods
```

## Updating Applications

Updating an application involves updating the image version.

### 1. Update the Deployment:

```
kubectl set image deployment/nginx-deployment nginx=nginx:1.16.1
```

### 2. Verify the Update:

```
kubectl rollout status deployment/nginx-deployment  
kubectl get pods
```

## ConfigMaps and Secrets

ConfigMaps

ConfigMaps allow you to decouple environment-specific configuration from your container images.

### 1. Create a ConfigMap:

```
kubectl create configmap example-config --from-literal=key1=value1 --  
from-literal=key2=value2
```

### 2. Use ConfigMap in a Pod:

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: configmap-pod  
spec:  
  containers:  
    - name: mycontainer  
      image: busybox  
      command: [ "sh", "-c", "env && sleep 3600" ]  
      env:
```



```

- name: SPECIAL_KEY
  valueFrom:
    configMapKeyRef:
      name: example-config
      key: key1

```

### 3. Apply the Pod configuration:

```
kubectl apply -f pod-configmap.yaml
```

### 4. Verify the Pod:

```
kubectl exec -it configmap-pod -- env
```

## Secrets

Secrets are used to store sensitive information, such as passwords, OAuth tokens, and ssh keys.

### 1. Create a Secret:

```
kubectl create secret generic db-user-pass --from-literal=username=admin --from-literal=password='S3cr3t!'
```

### 2. Use Secret in a Pod:

```

apiVersion: v1
kind: Pod
metadata:
  name: secret-pod
spec:
  containers:
  - name: mycontainer
    image: busybox
    command: [ "sh", "-c", "env && sleep 3600" ]
    env:
    - name: DB_USERNAME
      valueFrom:
        secretKeyRef:
          name: db-user-pass
          key: username
    - name: DB_PASSWORD
      valueFrom:
        secretKeyRef:
          name: db-user-pass
          key: password

```

### 3. Apply the Pod configuration:

```
kubectl apply -f pod-secret.yaml
```

### 4. Verify the Pod:

```
kubectl exec -it secret-pod -- env
```

## Persistent Storage

Kubernetes supports different types of persistent storage. Here's how to use a PersistentVolume and PersistentVolumeClaim.

### 1. Create a PersistentVolume (pv.yaml):

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-volume
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

### 2. Create a PersistentVolumeClaim (pvc.yaml):

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

### 3. Apply the PersistentVolume and PersistentVolumeClaim:

```
kubectl apply -f pv.yaml
kubectl apply -f pvc.yaml
```

### 4. Use PersistentVolumeClaim in a Pod:

```
apiVersion: v1
kind: Pod
metadata:
  name: pv-pod
spec:
  containers:
    - name: mycontainer
      image: busybox
      command: [ "sh", "-c", "echo 'Hello World' > /mnt/data/hello.txt
&& sleep 3600" ]
      volumeMounts:
        - mountPath: "/mnt/data"
          name: mypvc
  volumes:
    - name: mypvc
      persistentVolumeClaim:
        claimName: pv-claim
```

### 5. Apply the Pod configuration:

```
kubectl apply -f pod-pv.yaml
```

## 6. Verify the Pod:

```
kubectl exec -it pv-pod -- cat /mnt/data/hello.txt
```

## Networking in Kubernetes

Kubernetes networking allows communication between Pods, services, and external resources.

Network Policies

Network Policies allow you to control the traffic flow between Pods.

### 1. Create a Network Policy (network-policy.yaml):

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-nginx
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
      - podSelector:
          matchLabels:
            access: "true"
```

### 2. Apply the Network Policy:

```
kubectl apply -f network-policy.yaml
```

### 3. Verify the Network Policy:

```
kubectl get networkpolicies
```

## Monitoring and Logging

Monitoring and logging are crucial for maintaining the health and performance of your applications.

Prometheus and Grafana

### 1. Install Prometheus and Grafana using Helm:

```
helm repo add prometheus-community https://prometheus-
community.github.io/helm-charts
helm repo update
helm install prometheus prometheus-community/prometheus
helm install grafana prometheus-community/grafana
```

## 2. Access Grafana:

```
kubectl get svc --namespace default -w grafana
```

## 3. Set up Prometheus as a data source in Grafana and create dashboards to visualize metrics.

ELK Stack (Elasticsearch, Logstash, Kibana)

## 1. Install Elasticsearch and Kibana using Helm:

```
helm repo add elastic https://helm.elastic.co
helm repo update
helm install elasticsearch elastic/elasticsearch
helm install kibana elastic/kibana
```

## 2. Configure Logstash to collect logs from your applications and send them to Elasticsearch.

## 3. Access Kibana:

```
kubectl get svc --namespace default -w kibana
```

## 4. Use Kibana to visualize and analyze logs.

## Helm: Kubernetes Package Manager

Helm simplifies the deployment and management of applications on Kubernetes.

## 1. Install Helm:

```
curl https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3 | bash
```

## 2. Add a Helm Repository:

```
helm repo add stable https://charts.helm.sh/stable
helm repo update
```

## 3. Install an Application:

```
helm install my-release stable/nginx
```

## 4. Upgrade an Application:

```
helm upgrade my-release stable/nginx
```

## 5. Uninstall an Application:

```
helm uninstall my-release
```

## CI/CD with Kubernetes

Integrating CI/CD with Kubernetes enables automated deployments and continuous delivery of applications.

Jenkins

### 1. Install Jenkins using Helm:

```
helm repo add jenkins https://charts.jenkins.io
helm repo update
helm install jenkins jenkins/jenkins
```

### 2. Access Jenkins:

```
kubectl get svc --namespace default -w jenkins
```

### 3. Configure Jenkins to build and deploy your applications to Kubernetes.

GitLab CI/CD

### 1. Create a `.gitlab-ci.yml` file:

```
stages:
  - build
  - deploy

build:
  stage: build
  script:
    - docker build -t my-app:$CI_COMMIT_SHA .

deploy:
  stage: deploy
  script:
    - kubectl apply -f deployment.yaml
```

### 2. Push the configuration to your GitLab repository and set up runners to execute the CI/CD pipeline.

## Best Practices for Kubernetes

1. **Namespace Isolation:** Use namespaces to isolate environments (e.g., dev, test, prod).
2. **Resource Requests and Limits:** Define resource requests and limits for your Pods to ensure efficient resource utilization.
3. **Readiness and Liveness Probes:** Use readiness and liveness probes to monitor the health of your applications.
4. **Auto-scaling:** Use Horizontal Pod Autoscaler (HPA) to automatically scale your applications based on metrics.

5. **Security:** Implement Role-Based Access Control (RBAC) and network policies to secure your cluster.
6. **Backup and Disaster Recovery:** Regularly back up your etcd database and implement disaster recovery procedures.
7. **Logging and Monitoring:** Ensure comprehensive logging and monitoring to maintain the health and performance of your applications.

## Conclusion

Kubernetes is a powerful platform for managing containerized applications. This guide has provided a comprehensive overview of Kubernetes, from setting up a cluster to deploying applications, scaling, updating, and managing configurations and secrets. By following these hands-on examples and best practices, you can effectively utilize Kubernetes to streamline your development and operations workflows.